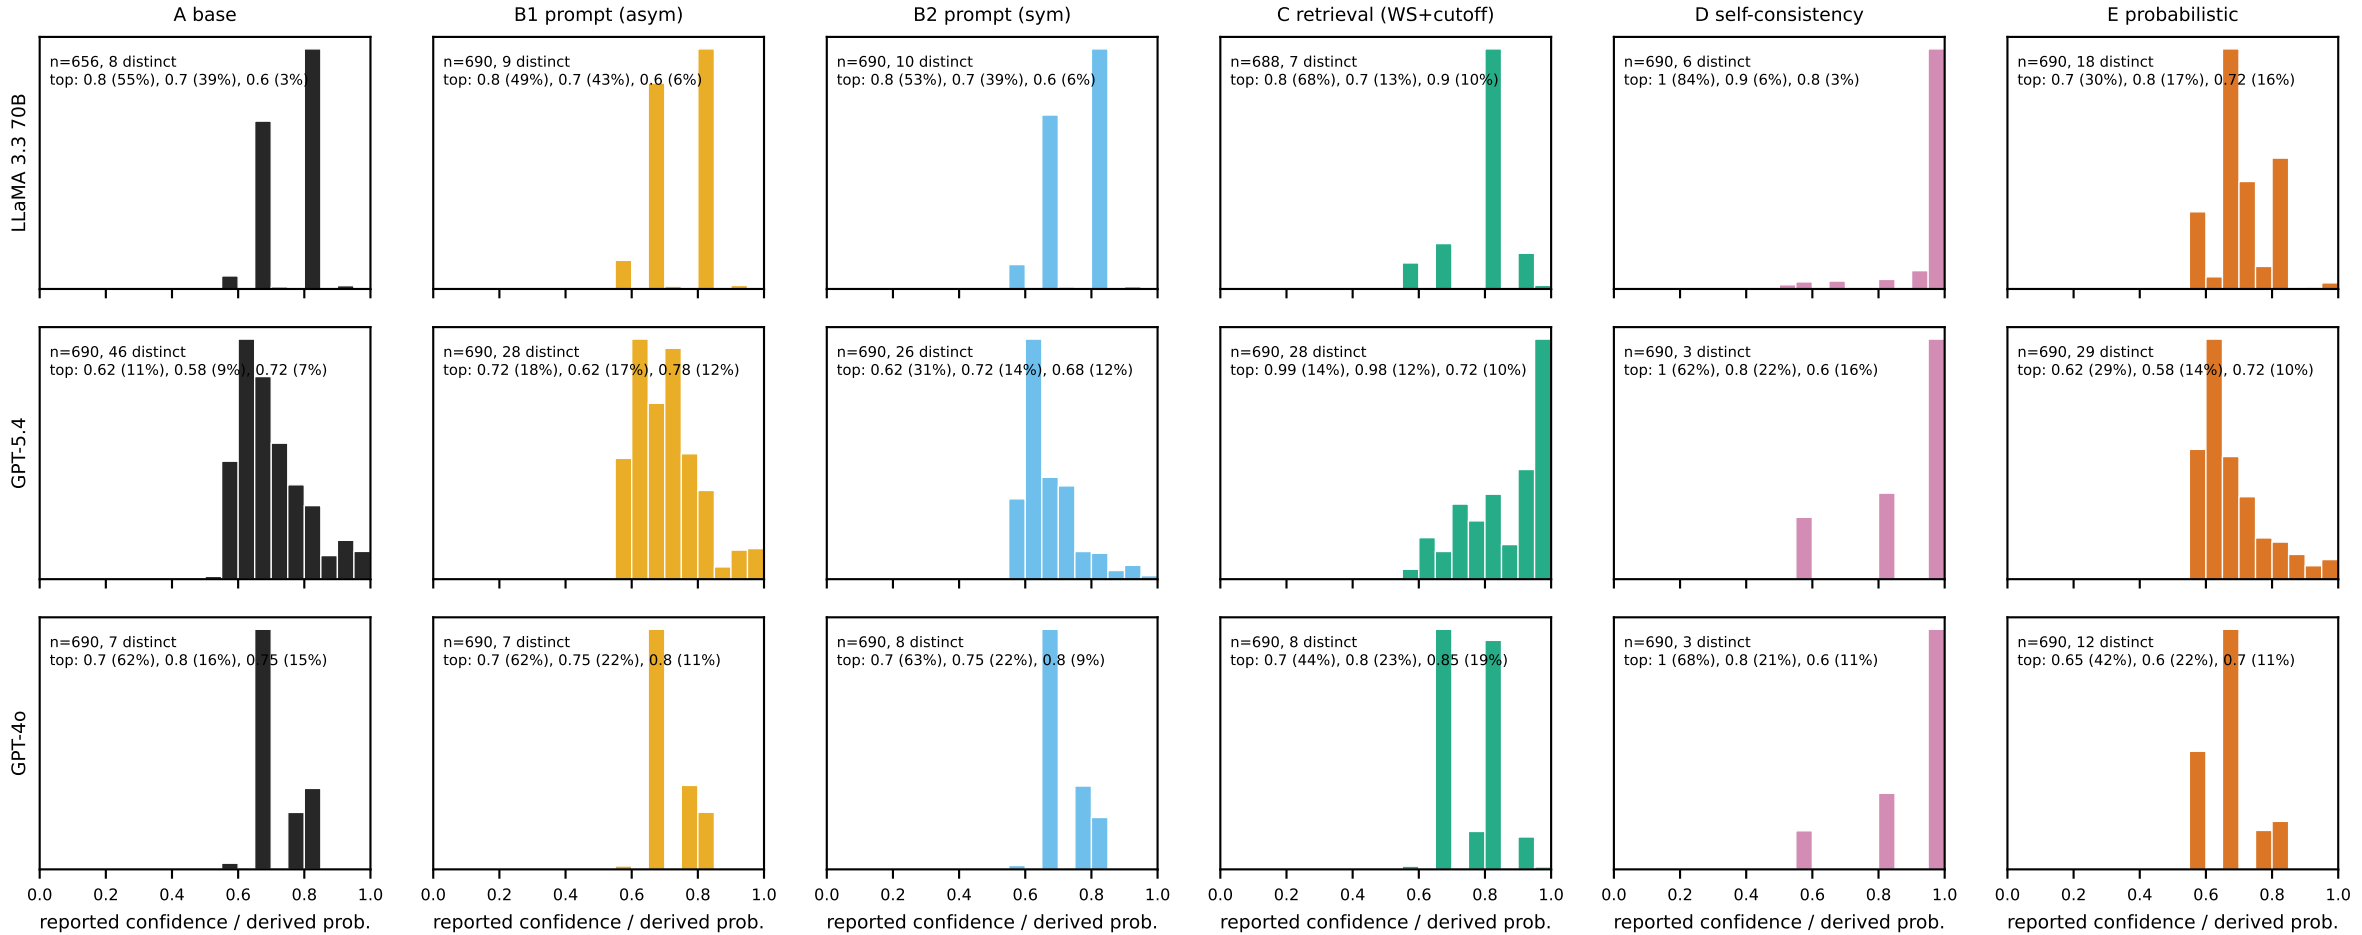
binary_merged: distribution of reported confidence per model x arm (0.05-wide bins; partition: all)
A base
B1 prompt (asym)
B2 prompt (sym)
C retrieval (WS+cutoff)
D self-consistency
E probabilistic
LLaMA 3.3 70B
GPT-5.4
GPT-4o
n=656, 8 distinct
top: 0.8 (55%), 0.7 (39%), 0.6 (3%)
n=690, 9 distinct
top: 0.8 (49%), 0.7 (43%), 0.6 (6%)
n=690, 10 distinct
top: 0.8 (53%), 0.7 (39%), 0.6 (6%)
n=688, 7 distinct
top: 0.8 (68%), 0.7 (13%), 0.9 (10%)
n=690, 6 distinct
top: 1 (84%), 0.9 (6%), 0.8 (3%)
n=690, 18 distinct
top: 0.7 (30%), 0.8 (17%), 0.72 (16%)
n=690, 46 distinct
top: 0.62 (11%), 0.58 (9%), 0.72 (7%)
n=690, 28 distinct
top: 0.72 (18%), 0.62 (17%), 0.78 (12%)
n=690, 26 distinct
top: 0.62 (31%), 0.72 (14%), 0.68 (12%)
n=690, 28 distinct
top: 0.99 (14%), 0.98 (12%), 0.72 (10%)
n=690, 3 distinct
top: 1 (62%), 0.8 (22%), 0.6 (16%)
n=690, 29 distinct
top: 0.62 (29%), 0.58 (14%), 0.72 (10%)
n=690, 7 distinct
top: 0.7 (62%), 0.8 (16%), 0.75 (15%)
n=690, 7 distinct
top: 0.7 (62%), 0.75 (22%), 0.8 (11%)
n=690, 8 distinct
top: 0.7 (63%), 0.75 (22%), 0.8 (9%)
n=690, 8 distinct
top: 0.7 (44%), 0.8 (23%), 0.85 (19%)
n=690, 3 distinct
top: 1 (68%), 0.8 (21%), 0.6 (11%)
n=690, 12 distinct
top: 0.65 (42%), 0.6 (22%), 0.7 (11%)
0.0
0.2
0.4
0.6
0.8
1.0
reported confidence / derived prob.