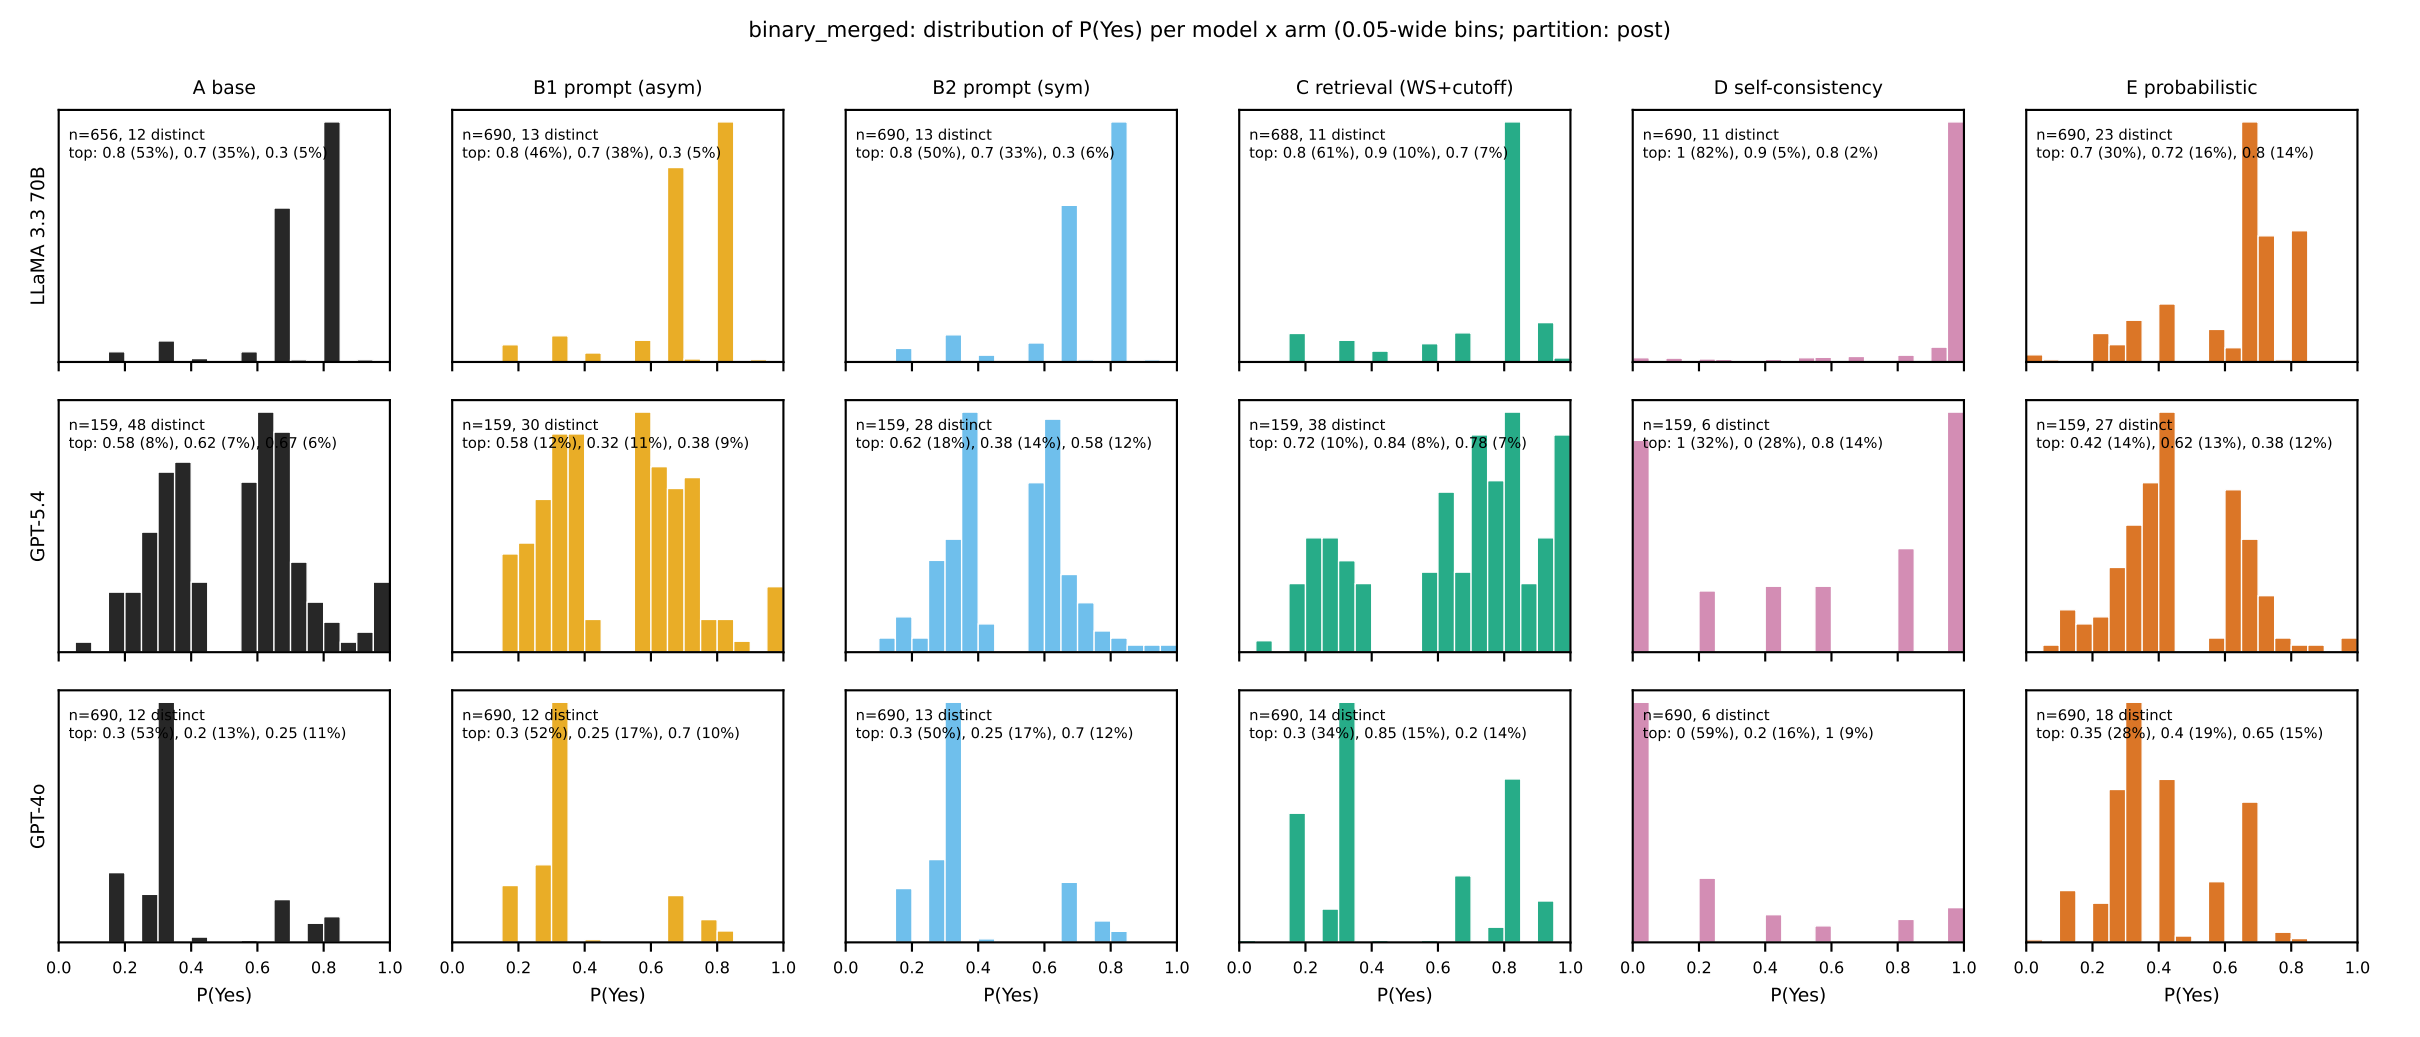
binary_merged: distribution of P(Yes) per model x arm (0.05-wide bins; partition: post)
A base
B1 prompt (asym)
B2 prompt (sym)
C retrieval (WS+cutoff)
D self-consistency
E probabilistic
LLaMA 3.3 70B
GPT-5.4
GPT-4o
n=656, 12 distinct
top: 0.8 (53%), 0.7 (35%), 0.3 (5%)
n=690, 13 distinct
top: 0.8 (46%), 0.7 (38%), 0.3 (5%)
n=690, 13 distinct
top: 0.8 (50%), 0.7 (33%), 0.3 (6%)
n=688, 11 distinct
top: 0.8 (61%), 0.9 (10%), 0.7 (7%)
n=690, 11 distinct
top: 1 (82%), 0.9 (5%), 0.8 (2%)
n=690, 23 distinct
top: 0.7 (30%), 0.72 (16%), 0.8 (14%)
n=159, 48 distinct
top: 0.58 (8%), 0.62 (7%), 0.67 (6%)
n=159, 30 distinct
top: 0.58 (12%), 0.32 (11%), 0.38 (9%)
n=159, 28 distinct
top: 0.62 (18%), 0.38 (14%), 0.58 (12%)
n=159, 38 distinct
top: 0.72 (10%), 0.84 (8%), 0.78 (7%)
n=159, 6 distinct
top: 1 (32%), 0 (28%), 0.8 (14%)
n=159, 27 distinct
top: 0.42 (14%), 0.62 (13%), 0.38 (12%)
n=690, 12 distinct
top: 0.3 (53%), 0.2 (13%), 0.25 (11%)
n=690, 12 distinct
top: 0.3 (52%), 0.25 (17%), 0.7 (10%)
n=690, 13 distinct
top: 0.3 (50%), 0.25 (17%), 0.7 (12%)
n=690, 14 distinct
top: 0.3 (34%), 0.85 (15%), 0.2 (14%)
n=690, 6 distinct
top: 0 (59%), 0.2 (16%), 1 (9%)
n=690, 18 distinct
top: 0.35 (28%), 0.4 (19%), 0.65 (15%)
0.0
0.2
0.4
0.6
0.8
1.0
P(Yes)