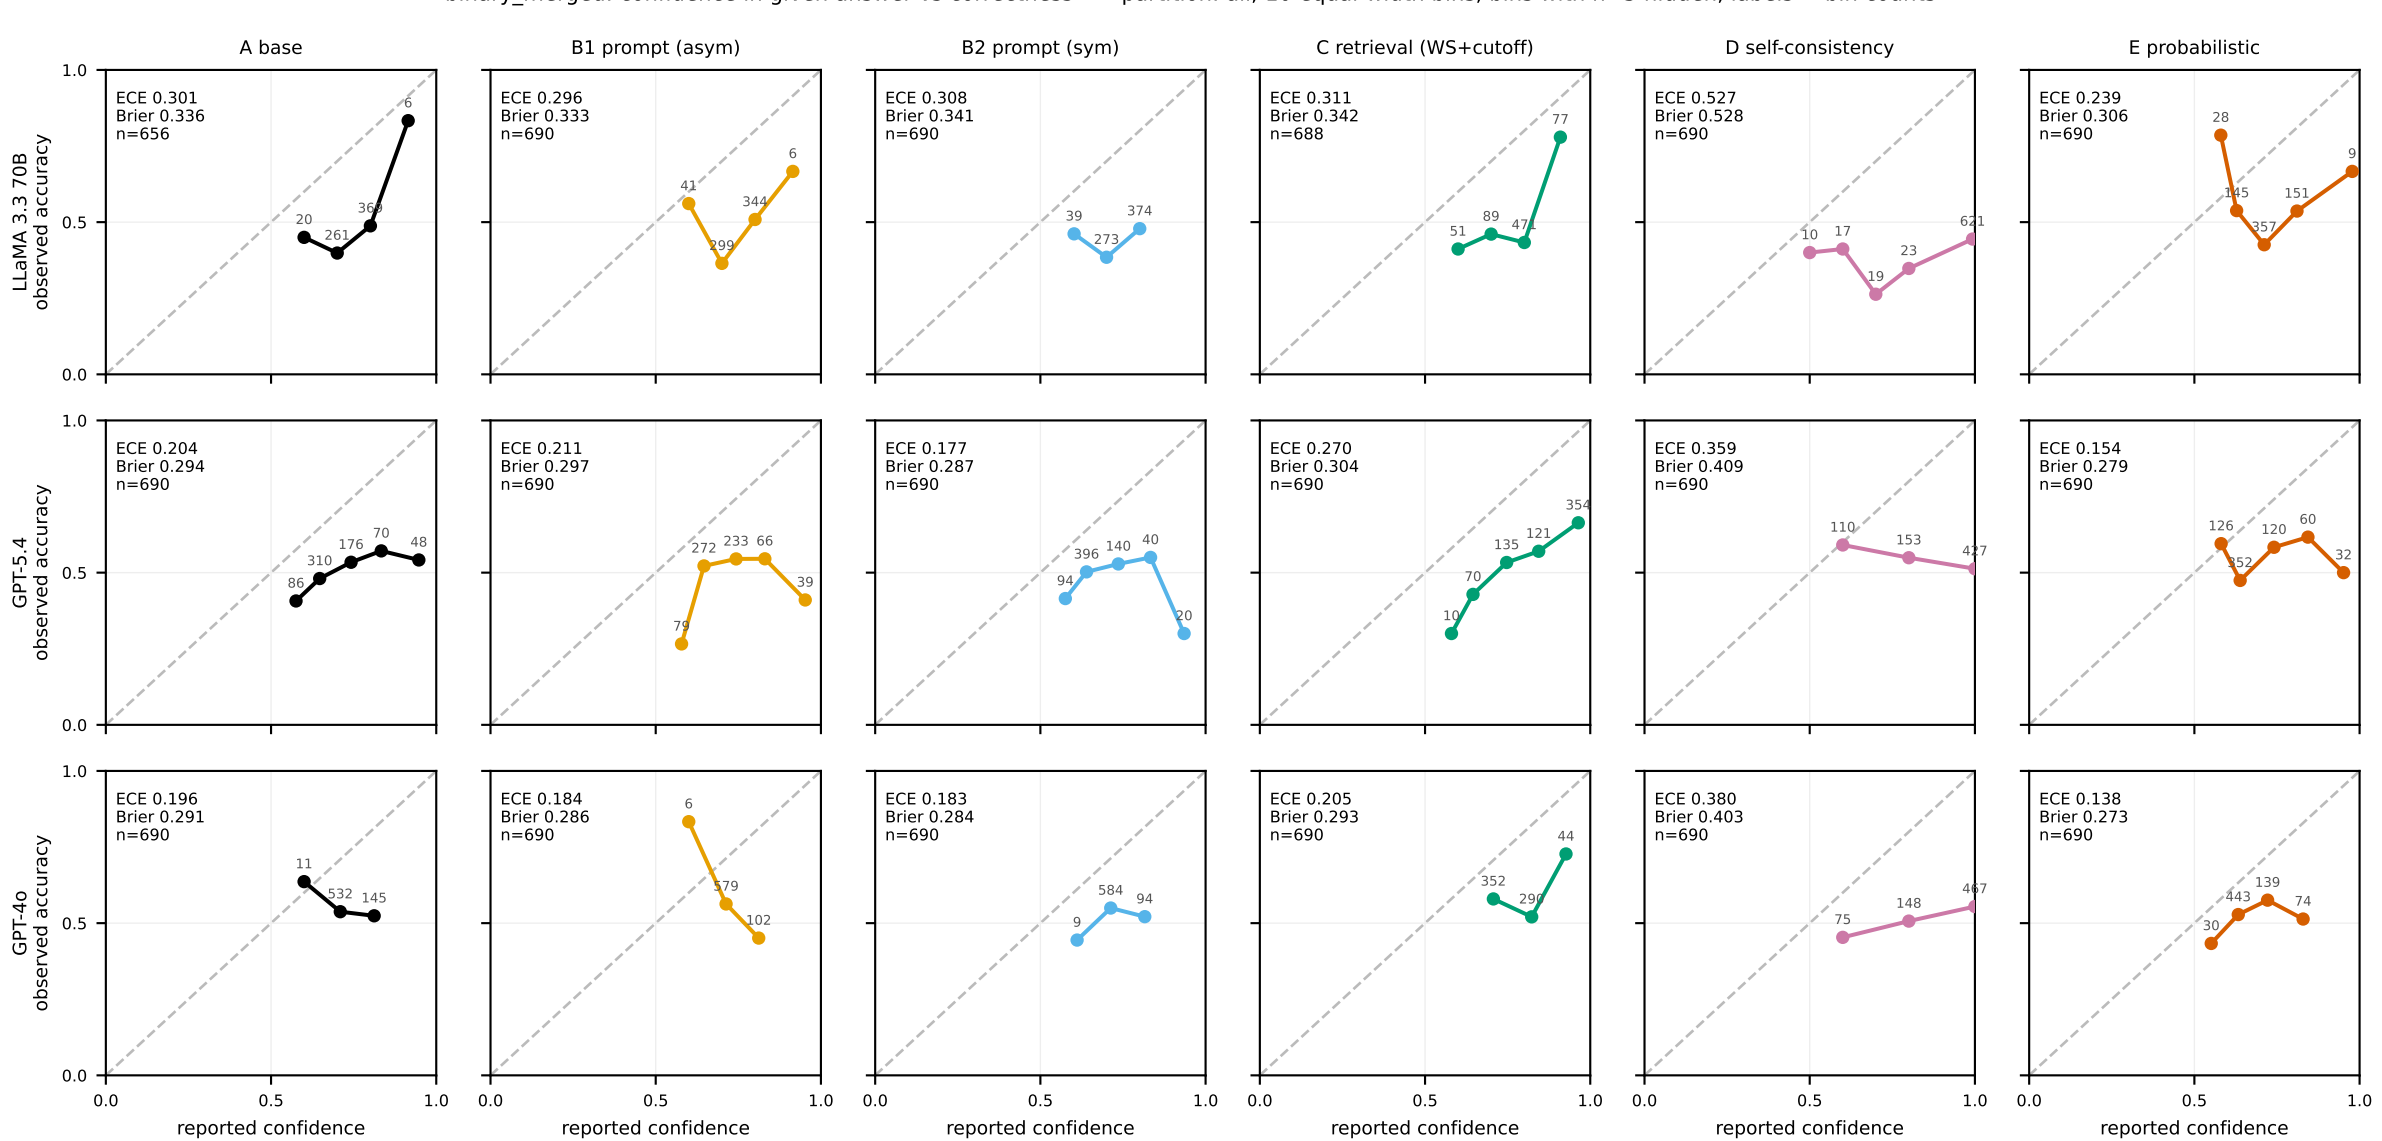
binary_merged: confidence in given answer vs correctness — partition: all; 10 equal-width bins, bins with n<5 hidden, labels = bin counts
A base
B1 prompt (asym)
B2 prompt (sym)
C retrieval (WS+cutoff)
D self-consistency
E probabilistic
LLaMA 3.3 70B observed accuracy
GPT-5.4 observed accuracy
GPT-4o observed accuracy
ECE 0.301
Brier 0.336
n=656
ECE 0.296
Brier 0.333
n=690
ECE 0.308
Brier 0.341
n=690
ECE 0.311
Brier 0.342
n=688
ECE 0.527
Brier 0.528
n=690
ECE 0.239
Brier 0.306
n=690
ECE 0.204
Brier 0.294
n=690
ECE 0.211
Brier 0.297
n=690
ECE 0.177
Brier 0.287
n=690
ECE 0.270
Brier 0.304
n=690
ECE 0.359
Brier 0.409
n=690
ECE 0.154
Brier 0.279
n=690
ECE 0.196
Brier 0.291
n=690
ECE 0.184
Brier 0.286
n=690
ECE 0.183
Brier 0.284
n=690
ECE 0.205
Brier 0.293
n=690
ECE 0.380
Brier 0.403
n=690
ECE 0.138
Brier 0.273
n=690
0.0
0.5
1.0
reported confidence