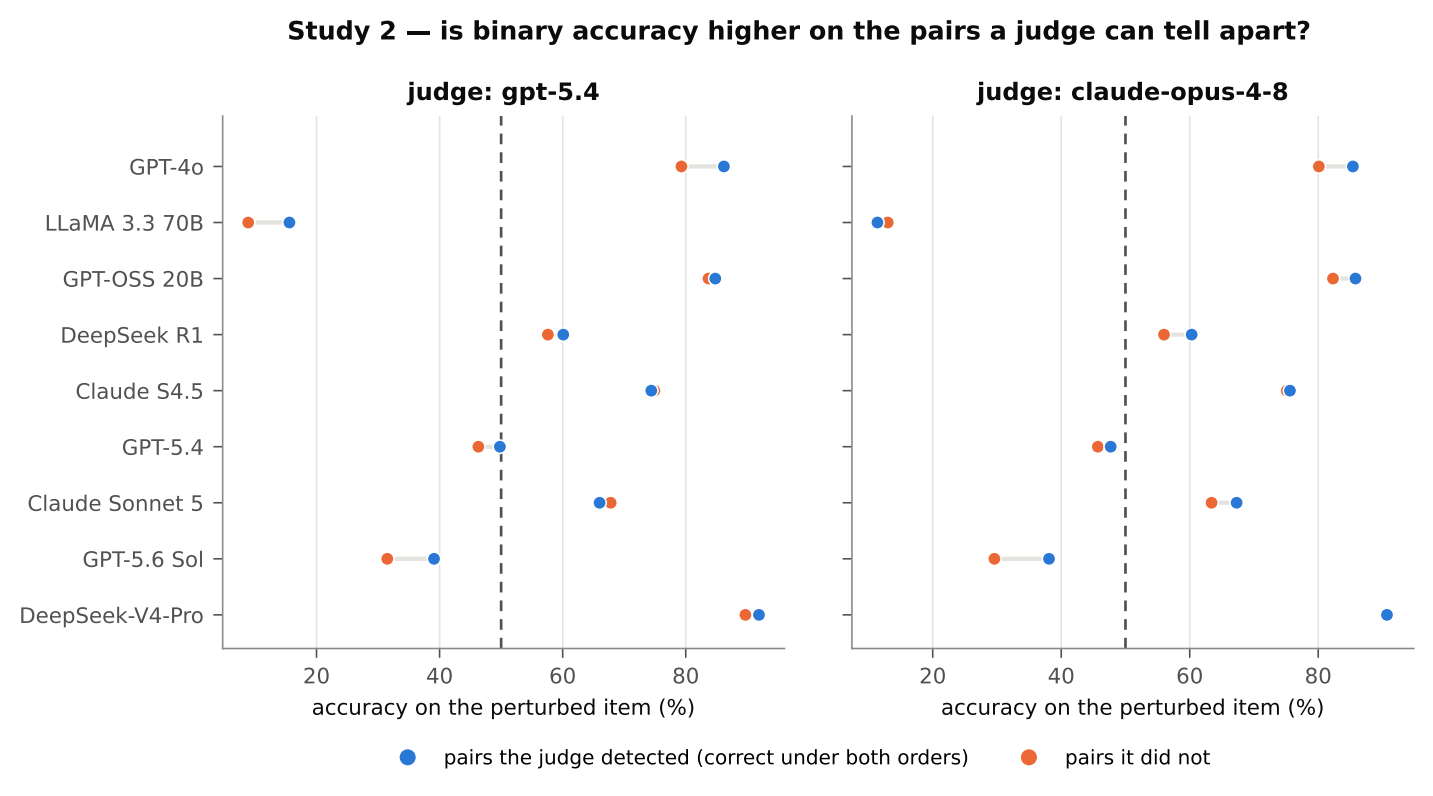
Study 2 — is binary accuracy higher on the pairs a judge can tell apart?
judge: gpt-5.4
judge: claude-opus-4-8
GPT-4o
LLaMA 3.3 70B
GPT-OSS 20B
DeepSeek R1
Claude S4.5
GPT-5.4
Claude Sonnet 5
GPT-5.6 Sol
DeepSeek-V4-Pro
20
40
60
80
accuracy on the perturbed item (%)
pairs the judge detected (correct under both orders)
pairs it did not